

Telecom Customer Churn Prediction



Project Overview

This project explores customer churn in a telecom company. The goal is to analyze customer behavior, identify patterns associated with churn, and build a model that predicts which customers are likely to leave.

Business Problem

As a stakeholder:

In a growing telecom company, I've observed a troubling pattern — we're losing customers at an increasing rate each month. Despite competitive pricing and a wide range of services, customer churn continues to rise, cutting into our recurring revenue and increasing customer acquisition costs. From our current customer data, out of 7,043 customers, 1,869 have churned — that's roughly 26.5% of our customer base. This is a significant red flag. After several internal reviews, it's clear that retaining existing customers is more cost-effective than acquiring new ones. But we currently lack a systematic approach to identify which customers are likely to leave — and why.

Objectives:

- Understand which factors most influence churn
- Build a model to predict the likelihood of churn
- Provide actionable recommendations to reduce churn

Dataset Description

Source: [Kaggle - Telco Customer Churn](#)

The dataset contains 21 columns including:

- Customer demographics
- Account information
- Services subscribed
- Monthly charges
- Whether they churned (Churn)

Load and Inspect the Data

In [713]:

```
#Importing Necessary libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

import warnings
warnings.filterwarnings('ignore')

sns.set(style="whitegrid", rc={
    'axes.grid': True,
    'axes.grid.axis': 'y',
```

```

'grid.color': 'dimgray',
'grid.linestyle': '-',
'grid.linewidth': 0.7
})
%matplotlib inline

```

In [714]:

```

# Load dataset
df = pd.read_csv("Data\WA_Fn-UseC_-Telco-Customer-Churn.csv")

print("Shape of dataset:", df.shape)
df.head()

```

Shape of dataset: (7043, 21)

Out[714]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No

5 rows x 21 columns



CLEANING AND EDA

CLEANING

In [717]:

```

# Checking for missing values and duplicates
print(df.isna().sum(), '\n')
print(f'Number of duplicated rows: {df.duplicated().sum()}')

```

```

customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0

```

Churn 0
dtype: int64

Number of duplicated rows: 0

In [718]:

```
df.describe()
```

Out[718]:

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

In [719]:

```
df.shape
```

Out[719]:

(7043, 21)

In [720]:

```
df.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
Column Non-Null Count Dtype
--- -
0 customerID 7043 non-null object
1 gender 7043 non-null object
2 SeniorCitizen 7043 non-null int64
3 Partner 7043 non-null object
4 Dependents 7043 non-null object
5 tenure 7043 non-null int64
6 PhoneService 7043 non-null object
7 MultipleLines 7043 non-null object
8 InternetService 7043 non-null object
9 OnlineSecurity 7043 non-null object
10 OnlineBackup 7043 non-null object
11 DeviceProtection 7043 non-null object
12 TechSupport 7043 non-null object
13 StreamingTV 7043 non-null object
14 StreamingMovies 7043 non-null object
15 Contract 7043 non-null object
16 PaperlessBilling 7043 non-null object
17 PaymentMethod 7043 non-null object
18 MonthlyCharges 7043 non-null float64
19 TotalCharges 7043 non-null object
20 Churn 7043 non-null object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

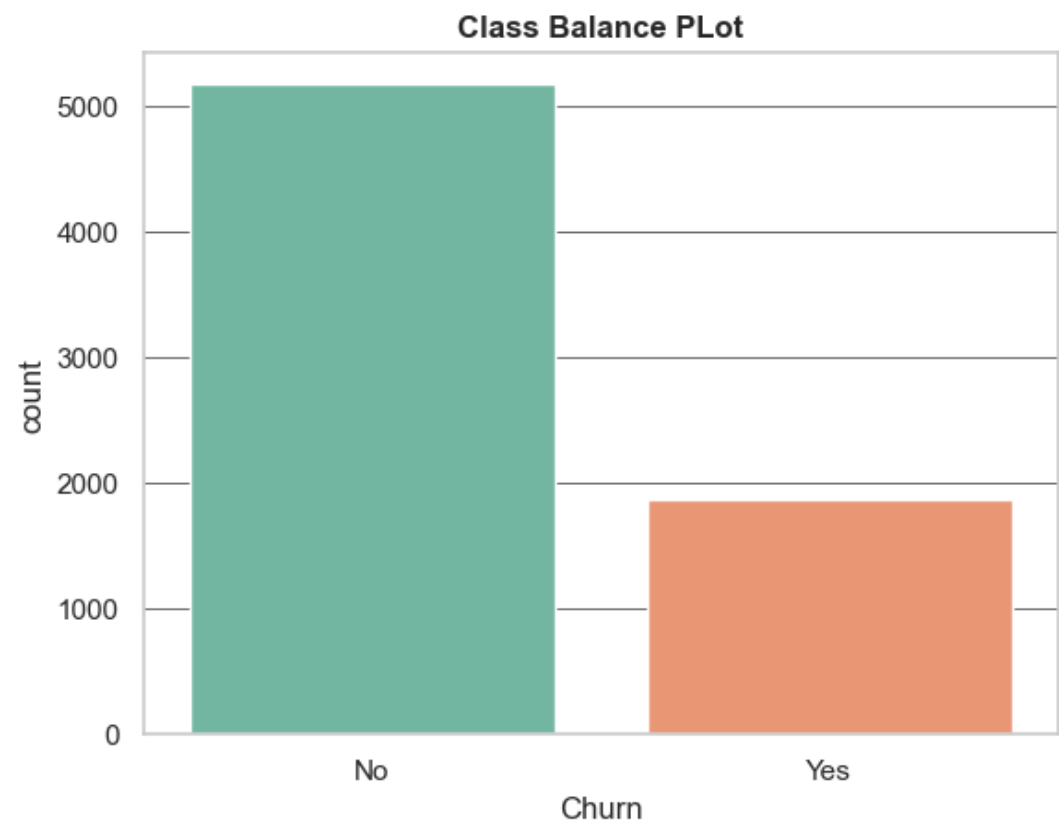
In [721]:

```
# Check the target class distribution  
df['Churn'].value_counts()
```

```
Out[721]:
Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

```
In [722]:
```

```
# Visualizing churn count
sns.countplot(x = 'Churn', data = df, hue = 'Churn', palette = 'Set2')
plt.title('Class Balance PLOT', weight = 'bold')
plt.show()
```



The plot above shows how our classes are distributed showing that the number of people who churn is less so we might need to use class balancing later on

- Now we can drop the data columns we will not use for this project

```
In [725]:
```

```
df = df.drop(columns = ['customerID'])
df.head()
```

```
Out[725]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBac
0	Female	0	Yes	No	1	No	No phone service	DSL	No	
1	Male	0	No	No	34	Yes	No	DSL	Yes	
2	Male	0	No	No	2	Yes	No	DSL	Yes	
3	Male	0	No	No	45	No	No phone service	DSL	Yes	
4	Female	0	No	No	0	No	No	DSL	No	

4	Female	0	No	No	2	Yes	No	Fiber optic	No	No
gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBa	

We can now divide the data into the customers who churn and those who did not to look for patterns in their own data sets.

Plotting them side by side or within the same plot will help us see how the data is different from the other.

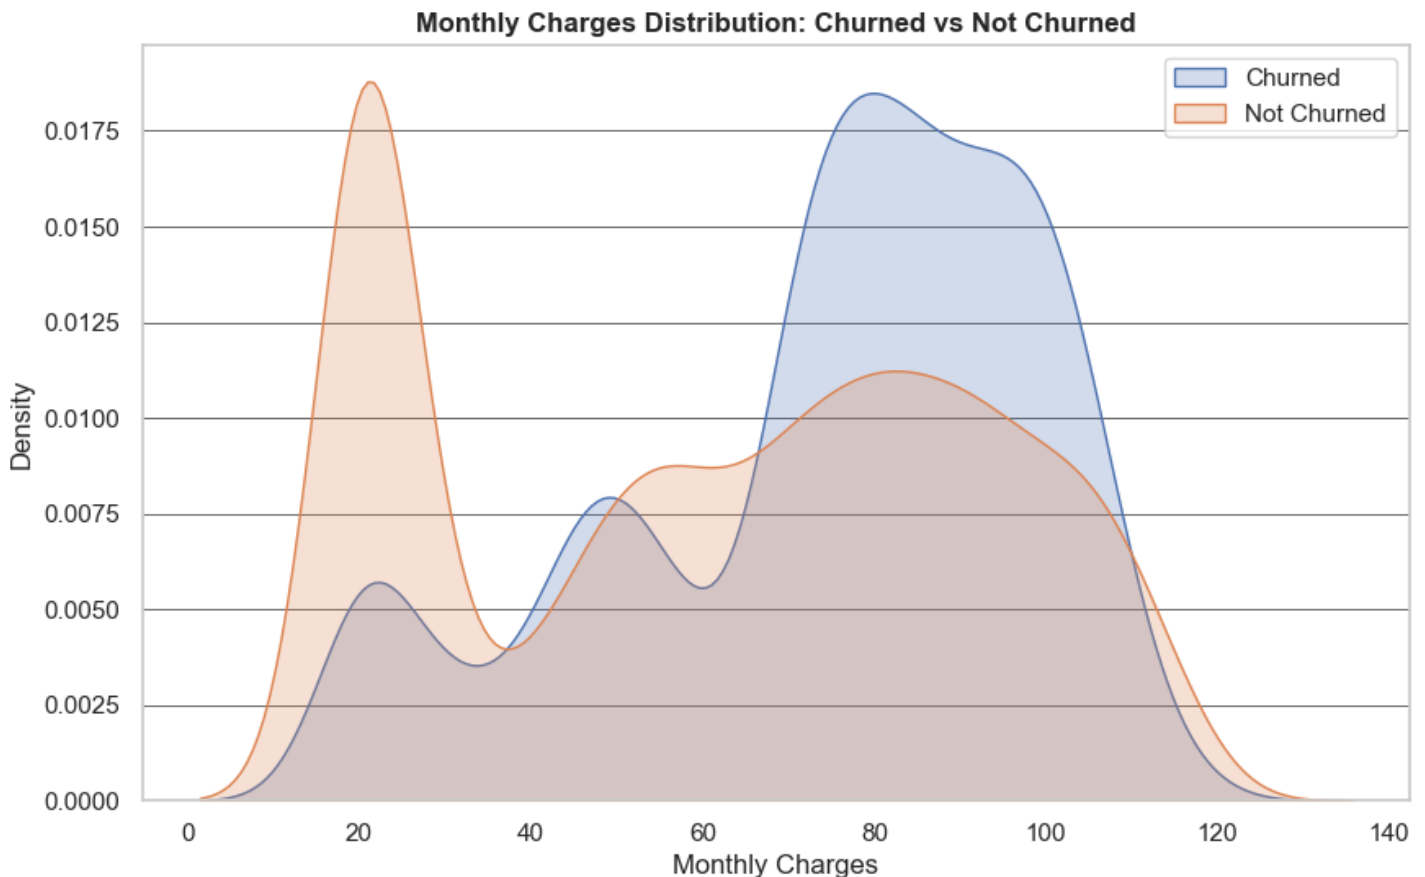
Looking at How Monthly Rates Affect Churning

In [728]:

```
#Splitting the data into two sets(Churned and not churned)
df_churned = df[df['Churn'] == 'Yes']
df_not_churned = df[df['Churn'] == 'No']
```

In [729]:

```
plt.figure(figsize=(10,6))
sns.kdeplot(df_churned['MonthlyCharges'], label='Churned', shade=True)
sns.kdeplot(df_not_churned['MonthlyCharges'], label='Not Churned', shade=True)
plt.title('Monthly Charges Distribution: Churned vs Not Churned', weight = 'bold')
plt.xlabel('Monthly Charges')
plt.legend()
plt.savefig('images\Monthly Rates VS Churning.jpg', dpi=300, bbox_inches='tight')
plt.show()
```



The plot above shows a density curve plot of how monthly rates affect if a customer churns or not.

OBSERVATION

- Churned are left-Skewed
- Not Churned are right-skewed

CONCLUSION

Higher Monthly Charges are more likely to churn

- The peak for churned customers is between 70 and 100, where their density is much higher than that of non-churned customers
- This suggests that customers paying higher monthly charges are more likely to churn

Lower Monthly Charges are less likely to churn

- The not churned group (orange) shows a strong peak around \$20, a region where the churned group is relatively low.
- This indicates that customers paying low monthly charges tend to stay

Looking If Gender Affects Churning

In [733]:

```
# first we need to know if the class is balanced to get a better reading
df['gender'].value_counts()
```

Out[733]:

```
gender
Male      3555
Female    3488
Name: count, dtype: int64
```

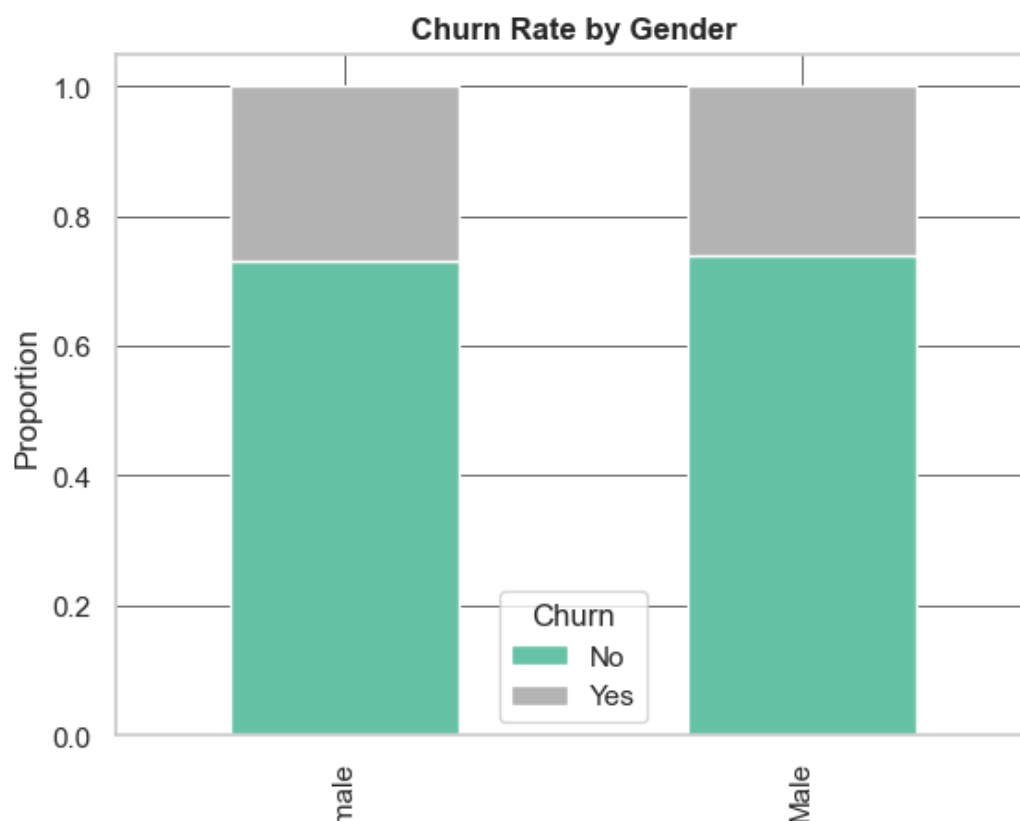
So the data is balanced now to look at how it affects churning

In [735]:

```
churn_by_gender = df.groupby('gender')['Churn'].value_counts(normalize=True).unstack()
print(churn_by_gender)
```

```
churn_by_gender.plot(kind='bar', stacked=True, colormap='Set2')
plt.title('Churn Rate by Gender', weight = 'bold')
plt.ylabel('Proportion')
plt.savefig('images\Gender VS Churning.jpg', dpi=300, bbox_inches='tight')
plt.show()
```

```
Churn      No      Yes
gender
Female    0.730791  0.269209
Male      0.738397  0.261603
```



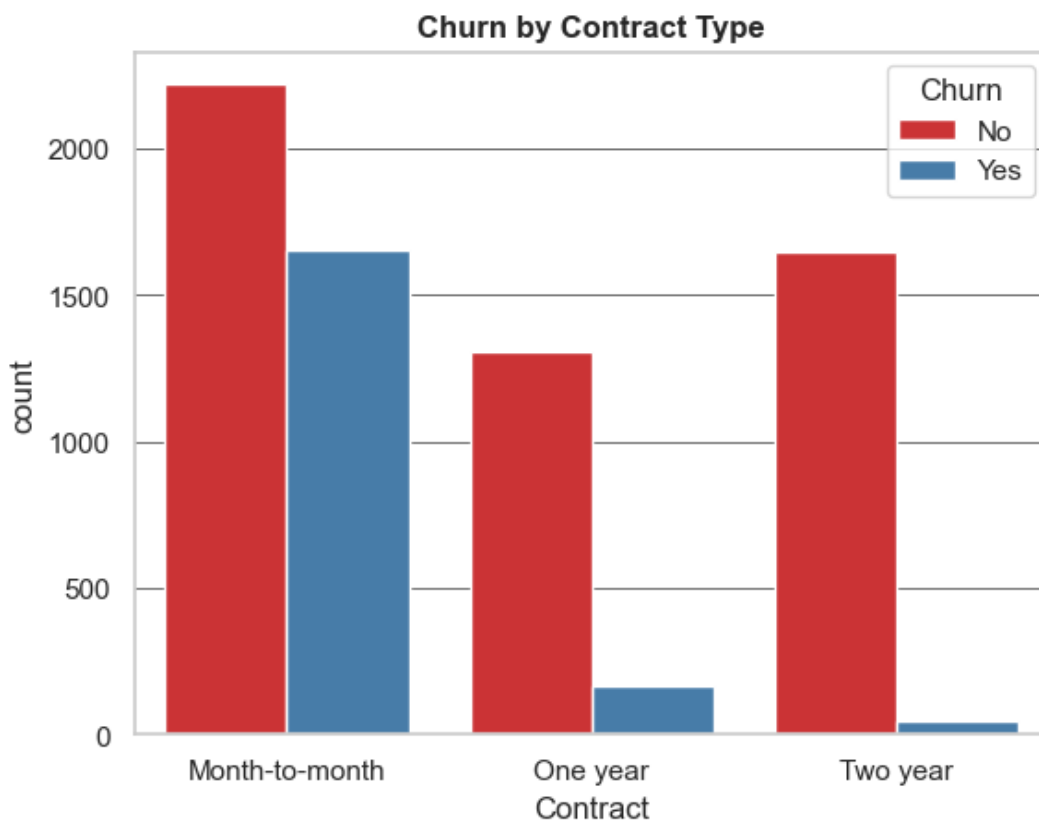
So gender does not carry too much based on the plot above since the number of churned is almost similar in both genders

Contract vs Churning

Here we will look at how contracts(month-month, one year, two years) affect churning

In [739]:

```
sns.countplot(x='Contract', hue='Churn', data=df, palette='Set1')
plt.title('Churn by Contract Type', weight = 'bold')
plt.xticks(rotation=0)
plt.savefig('images\Contract VS Churning.jpg', dpi=300, bbox_inches='tight')
plt.show()
```



The plot shows that most people who churn the services are the month-month customers

CONCLUSION

- getting customers to commit to longer contracts will reduce churning

RECOMENDATION

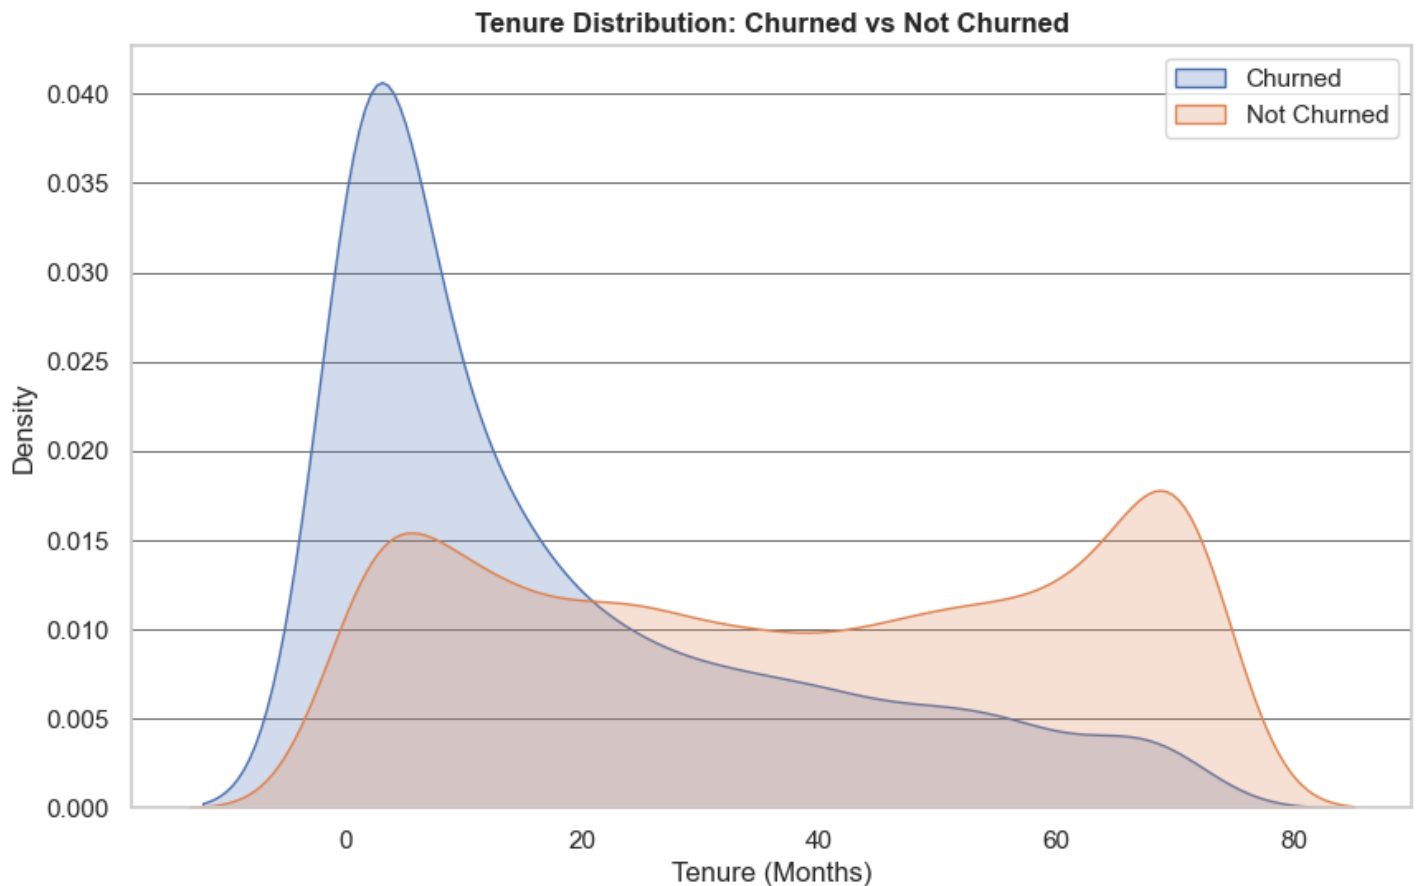
- Offer incentives (like discounts, perks, or exclusive services) to month-to-month customers to convert them to longer contracts and reduce churn.

Tenure(How Long the in Months a customer has stayed with the company) VS Churn

In [742]:

```
plt.figure(figsize=(10,6))
sns.kdeplot(df_churned['tenure'], label='Churned', shade=True)
sns.kdeplot(df_not_churned['tenure'], label='Not Churned', shade=True)
plt.title('Tenure Distribution: Churned vs Not Churned', weight = 'bold')
```

```
plt.xlabel('Tenure (Months)')
plt.legend()
plt.savefig('images\Tenure VS Churning.jpg', dpi=300, bbox_inches='tight')
plt.show()
```



This plot above shows a density curve of how tenure is distributed between churned and not churned customers

The observation and conclusion is that customers that have less tenure are more likely to churn the services while the customers who don't have a rather long tenure with the company

Recommendation : We try and get the customer to commit for about 6-12 months to reduce the chance of churning

Churning VS Internet Service types

In [745]:

```
#1st Look how this class is balanced
df['InternetService'].value_counts()
```

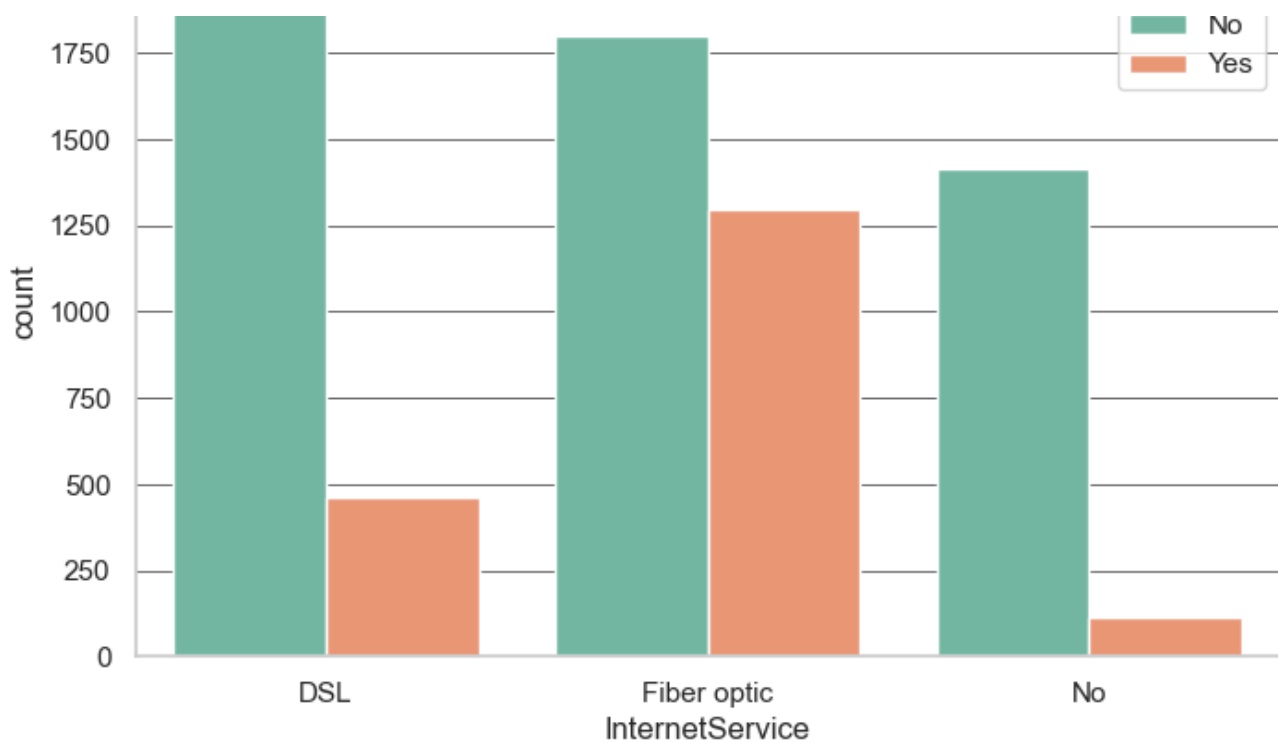
Out[745]:

```
InternetService
Fiber optic    3096
DSL            2421
No             1526
Name: count, dtype: int64
```

In [746]:

```
plt.figure(figsize=(8,5))
sns.countplot(x='InternetService', hue='Churn', data=df, palette='Set2')
plt.title('Churn by Internet Service Type', weight = 'bold')
plt.savefig('images\Internet Service Type VS Churning.jpg', dpi=300, bbox_inches='tight')
plt.show()
```





CONCLUSION

- Customers using Fiber optic internet are churning at higher rates compared to those using DSL or no internet service.

RECOMENDATION

- To reduce churn, consider offering incentives or easy options to switch internet service types, especially from fiber optic to other available options that might better fit customer needs or satisfaction

Churning VS Payment method

In [749]:

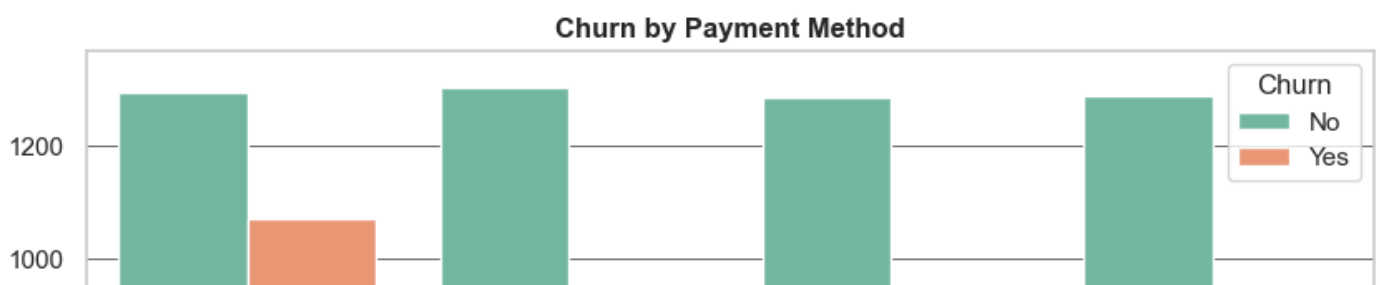
```
# Looking at the data in the Payment Method Column
df['PaymentMethod'].value_counts()
```

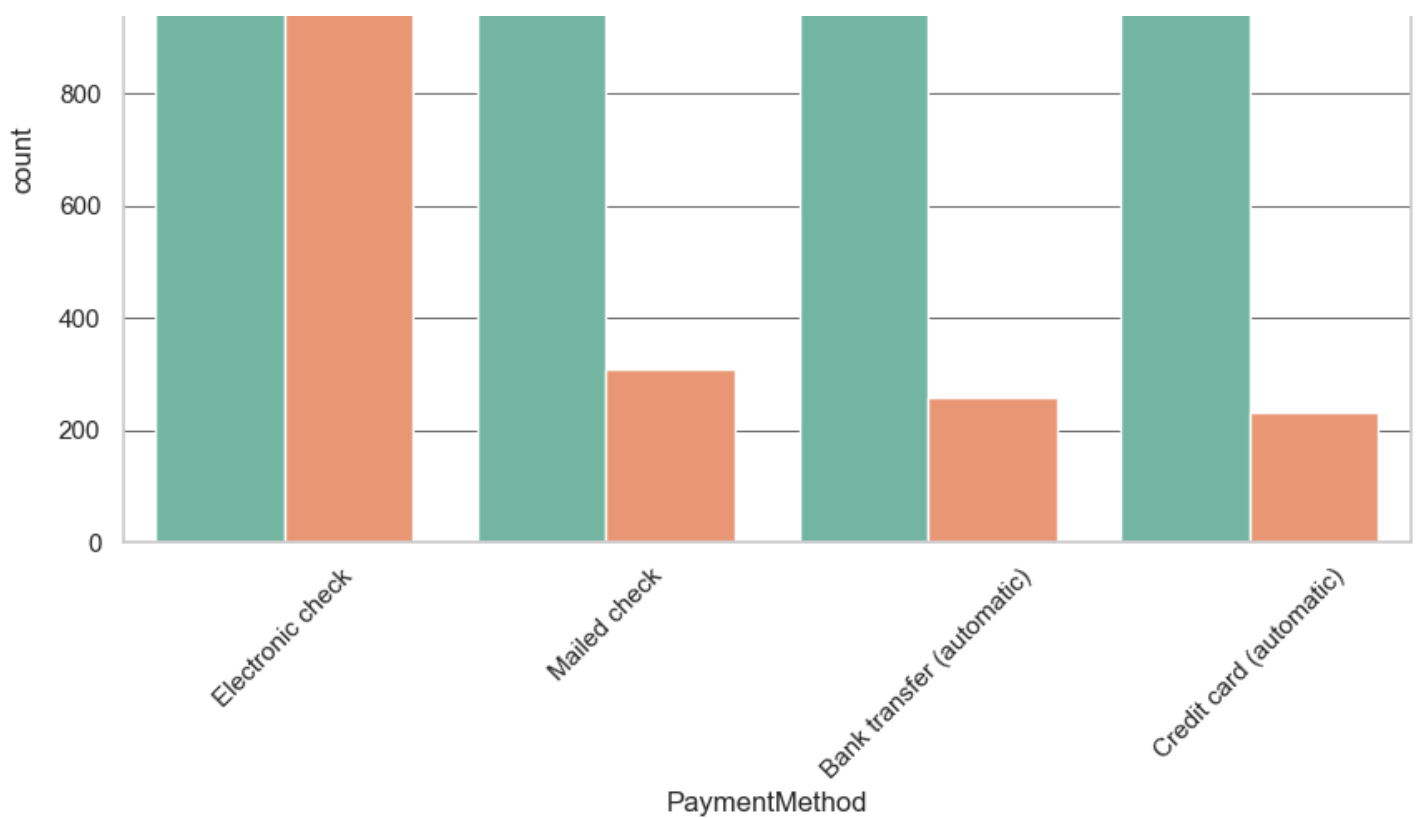
Out[749]:

```
PaymentMethod
Electronic check    2365
Mailed check       1612
Bank transfer (automatic)  1544
Credit card (automatic)  1522
Name: count, dtype: int64
```

In [750]:

```
plt.figure(figsize=(10,6))
sns.countplot(x='PaymentMethod', hue='Churn', data=df, palette='Set2')
plt.title('Churn by Payment Method', weight = 'bold')
plt.xticks(rotation=45)
plt.savefig('images\Payment VS Churning.jpg', dpi=300, bbox_inches='tight')
plt.show()
```





OBSERVATION

- Though the payment methods are roughly almost the same in terms of the number of customers who use them **Electronic check** has a much higher churn rate than the others

RECOMENDAION

- Investigate why customers who use **Electronic check** have a higher rate than the rest

SeniorCitizen VS Churn

In [753]:

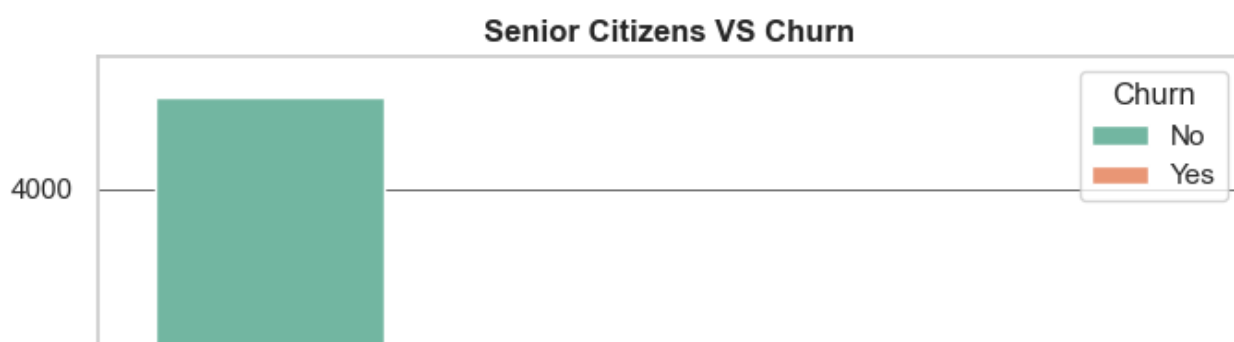
```
# Look at the distribution of Senior citizens columns
df['SeniorCitizen'].value_counts()
```

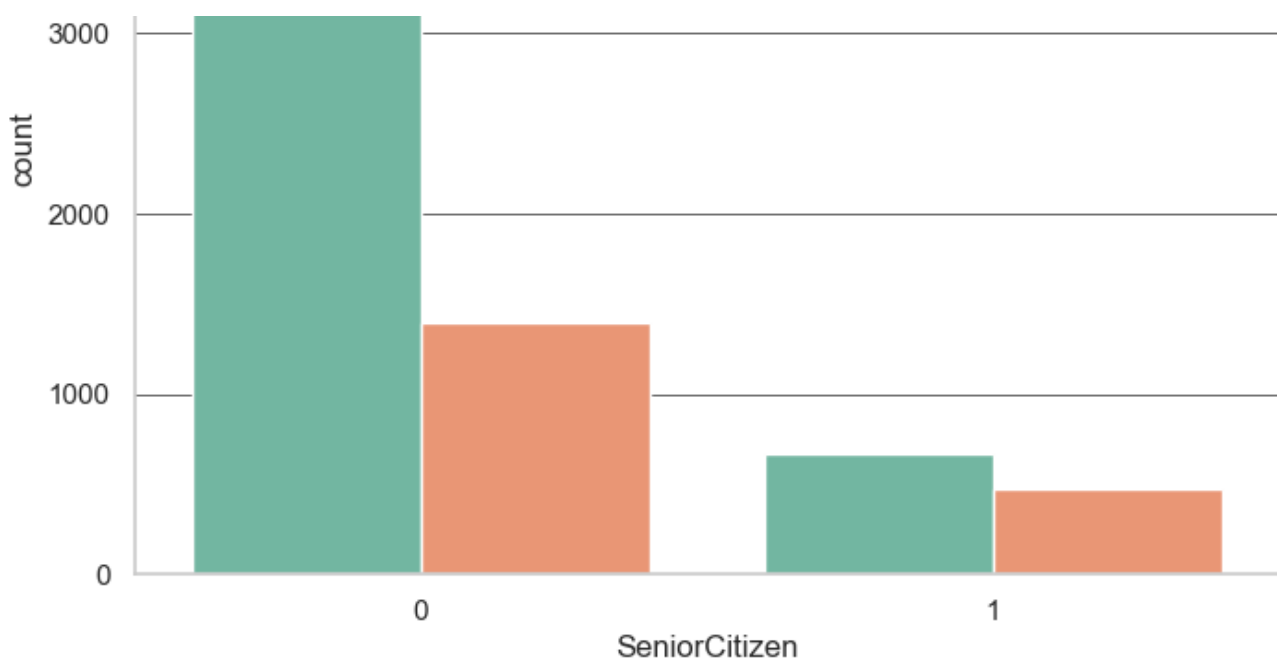
Out[753]:

```
SeniorCitizen
0      5901
1      1142
Name: count, dtype: int64
```

In [754]:

```
plt.figure(figsize=(8,6))
sns.countplot(x = 'SeniorCitizen', hue = 'Churn', data = df, palette = 'Set2')
plt.title('Senior Citizens VS Churn', weight = 'bold')
plt.savefig('images\Senior Citizens VS Churning.jpg', dpi=300, bbox_inches='tight')
plt.show()
```





OBSERVATION

- Based on the plot above the percentage of senior citizens who churn is higher than the the non-senior citizens

CONCLUSION

- Offer more serices that are more suitable to senior citizens to discarouge churning

MODELING

Data Preprocessing

1. Drop unnecessary columns

In [759]:

```
df = df.drop(columns = ['TotalCharges', 'gender', 'PhoneService'])
df.head()
```

Out[759]:

	SeniorCitizen	Partner	Dependents	tenure	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection
0	0	Yes	No	1	No phone service	DSL	No	Yes	No
1	0	No	No	34	No	DSL	Yes	No	Yes
2	0	No	No	2	No	DSL	Yes	Yes	No
3	0	No	No	45	No phone service	DSL	Yes	No	Yes
4	0	No	No	2	No	Fiber optic	No	No	No

In [760]:

```
df.columns
```

```
df.columns
```

Out[760]:

```
Index(['SeniorCitizen', 'Partner', 'Dependents', 'tenure', 'MultipleLines',  
      'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',  
      'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',  
      'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges', 'Churn'],  
      dtype='object')
```

2. Encode categorical variables

In [762]:

```
# Importing libraries  
from sklearn.preprocessing import OneHotEncoder  
  
df_encoded = pd.get_dummies(df, drop_first = True, dtype = int)  
df_encoded.head()
```

Out[762]:

	SeniorCitizen	tenure	MonthlyCharges	Partner_Yes	Dependents_Yes	MultipleLines_No phone service	MultipleLines_Yes	InternetService_
0	0	1	29.85	1	0	1	0	
1	0	34	56.95	0	0	0	0	
2	0	2	53.85	0	0	0	0	
3	0	45	42.30	0	0	1	0	
4	0	2	70.70	0	0	0	0	

5 rows x 28 columns



3. Split into training and test sets & Scale numeric features (for logistic regression)

In [764]:

```
# Importing Libraries  
from sklearn.preprocessing import StandardScaler  
from sklearn.model_selection import train_test_split  
#Define X and Y variables  
X = df_encoded.drop(columns = ['Churn_Yes'])  
y = df_encoded['Churn_Yes']  
  
# Split data into Train and Test  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state  
= 42, stratify=y)  
  
#Scale the X variables  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```

ML Modeling

Logistic Regression

In [767]:

```
# Importing Libraries  
from sklearn.linear_model import LogisticRegression  
  
#Initiate and fit the model
```

```
model0 = LogisticRegression()  
model0.fit(X_train_scaled, y_train)
```

Out[767]:

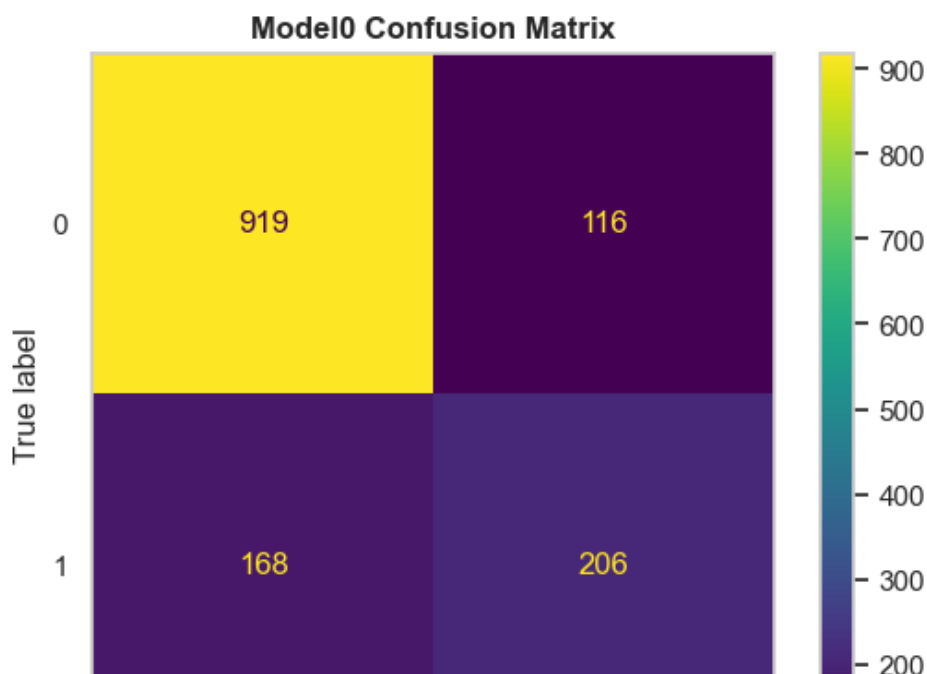
```
▼ LogisticRegression  
LogisticRegression()
```

Model Evaluation

In [769]:

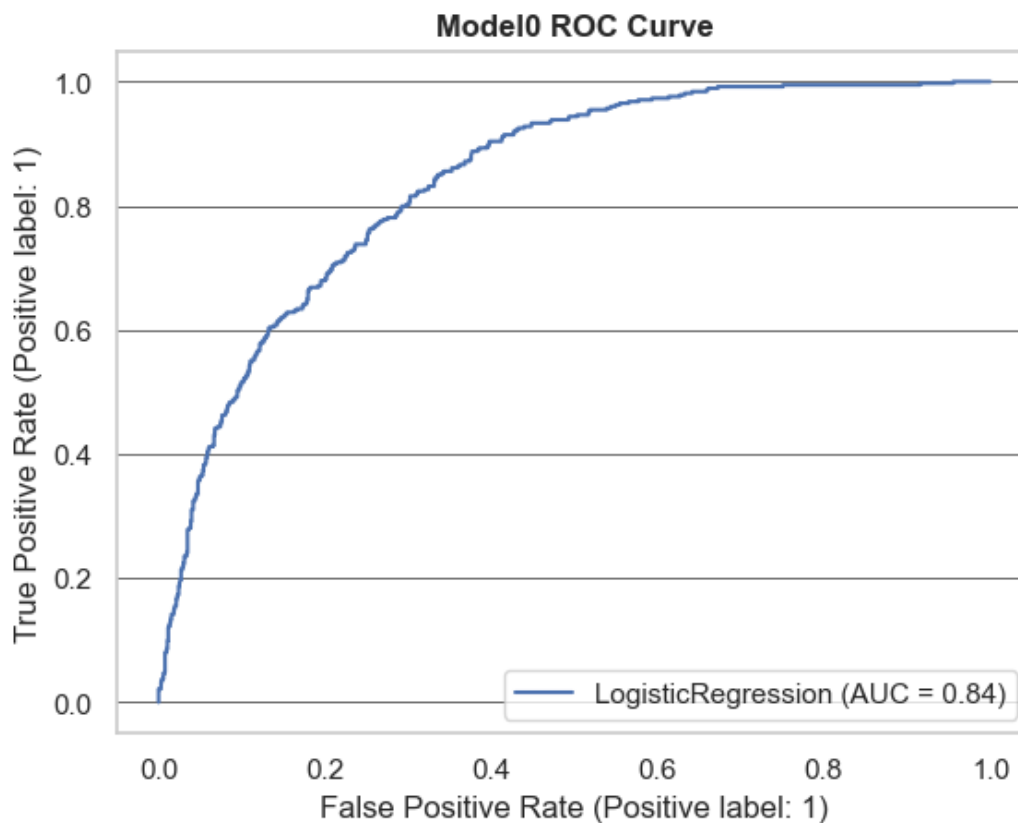
```
# importing libraries  
from sklearn.metrics import (classification_report, recall_score, f1_score, roc_auc_score  
, RocCurveDisplay,  
                             accuracy_score, precision_score, confusion_matrix, Confusion  
MatrixDisplay, roc_curve)  
  
# predicting the X_test_scaled  
y_pred = model0.predict(X_test_scaled)  
  
y_proba = model0.predict_proba(X_test_scaled)[:, 1]  
  
# Displaying the Confusion matrix  
cm = confusion_matrix(y_test, y_pred)  
matrix = ConfusionMatrixDisplay(confusion_matrix = cm)  
plt.figure(figsize = (8, 6))  
matrix.plot()  
plt.grid(False)  
plt.title('Model0 Confusion Matrix', weight = 'bold')  
plt.show()  
  
#ROC Curve Display  
plt.figure(figsize = (8, 6))  
RocCurveDisplay.from_estimator(model0, X_test_scaled, y_test)  
plt.title('Model0 ROC Curve', weight = 'bold')  
plt.show()  
  
# Get the scores for my model  
print(f'Model Report: {classification_report(y_test, y_pred)}')  
print(f'\nAccuracy Score: {accuracy_score(y_test, y_pred)}')  
print(f'\nROC_AUC Score: {roc_auc_score(y_test, y_proba)}')  
print(f'\nPrecision Score: {precision_score(y_test, y_pred)}')  
print(f'\nF1 Score: {f1_score(y_test, y_pred)}')  
print(f'\nRecall Score: {recall_score(y_test, y_pred)}')
```

<Figure size 800x600 with 0 Axes>





<Figure size 800x600 with 0 Axes>



Model Report:			precision	recall	f1-score	support
	0	0.85	0.89	0.87	1035	
	1	0.64	0.55	0.59	374	
	accuracy			0.80	1409	
	macro avg	0.74	0.72	0.73	1409	
	weighted avg	0.79	0.80	0.79	1409	

Accuracy Score: 0.7984386089425124

ROC_AUC Score: 0.8391123511328115

Precision Score: 0.639751552795031

F1 Score: 0.5919540229885057

Recall Score: 0.5508021390374331

- The scores above show that my model is doing well but could do better especially with precision score so for that I need to do feature selection to try and see which fetures have low impact on my model and try and do another model after dropping the features

Feature Selection

In [772]:

```
# Importing libraries
from sklearn.feature_selection import RFE

selector = RFE(LogisticRegression(), n_features_to_select=10)
selector.fit(X_train_scaled, y_train)
```

```
selected_features = X.columns[selector.support_]
print(selected_features)
```

```
Index(['tenure', 'MonthlyCharges', 'MultipleLines_Yes',
      'InternetService_Fiber optic', 'StreamingTV_No internet service',
      'StreamingTV_Yes', 'StreamingMovies_No internet service',
      'StreamingMovies_Yes', 'Contract_One year', 'Contract_Two year'],
      dtype='object')
```

From the code cell above the top 10 features with the biggest impact are:

- tenure
- MonthlyCharges
- MultipleLines_Yes
- InternetService_Fiber optic
- StreamingTV_No internet service
- StreamingTV_Yes
- StreamingMovies_No internet service
- StreamingMovies_Yes
- Contract_One year
- Contract_Two year

So I will build another model(model1) to compare how it does with the 1st model(model0)

In [775]:

```
# Building Modell
# Top 10 selected features
top_features = [
    'tenure',
    'MonthlyCharges',
    'MultipleLines_Yes',
    'InternetService_Fiber optic',
    'StreamingTV_No internet service',
    'StreamingTV_Yes',
    'StreamingMovies_No internet service',
    'StreamingMovies_Yes',
    'Contract_One year',
    'Contract_Two year'
]
# Variable initialization
X_top = df_encoded[top_features]
y = df_encoded['Churn_Yes']

# Split
X_train_top, X_test_top, y_train, y_test = train_test_split(X_top, y, test_size = 0.2, r
andom_state = 42, stratify = y)

# Scale
scaler1 = StandardScaler()
X_train_top_scaled = scaler1.fit_transform(X_train_top)
X_test_top_scaled = scaler1.transform(X_test_top)

# Build and fit Modell
modell1 = LogisticRegression(class_weight = 'balanced', random_state = 42)
modell1.fit(X_train_top_scaled, y_train)

# Predict y_test
y_pred2 = modell1.predict(X_test_top_scaled)

#Modell1 Probability Prediction
y_proba2 = modell1.predict_proba(X_test_top_scaled)[: , 1]

# Confusion Matrix display for Modell1
cm1 = confusion_matrix(y_test, y_pred2)
matrix1 = ConfusionMatrixDisplay(confusion_matrix = cm1)
plt.figure(figsize = (8, 6))
```

```

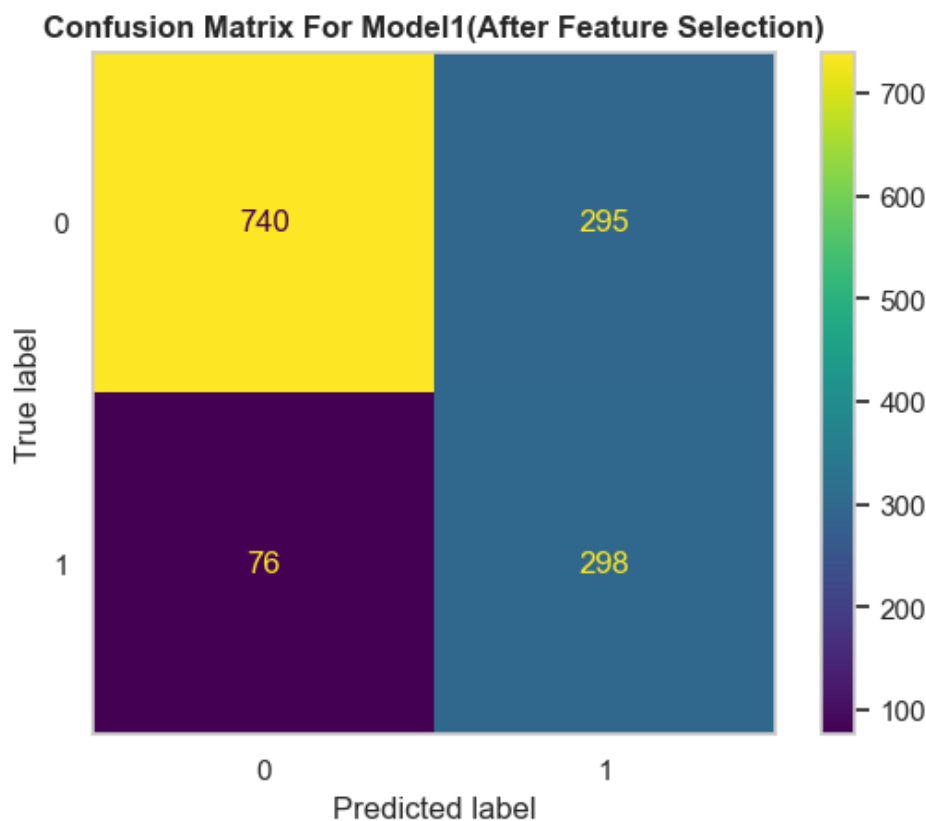
matrix1.plot()
plt.grid(False)
plt.title('Confusion Matrix For Modell(After Feature Selection)', weight = 'bold')
plt.show()

# ROC Curve for modell
plt.figure(figsize = (8, 6))
RocCurveDisplay.from_estimator(modell, X_test_top_scaled, y_test)
plt.title('Modell ROC Curve', weight = 'bold')
plt.show()

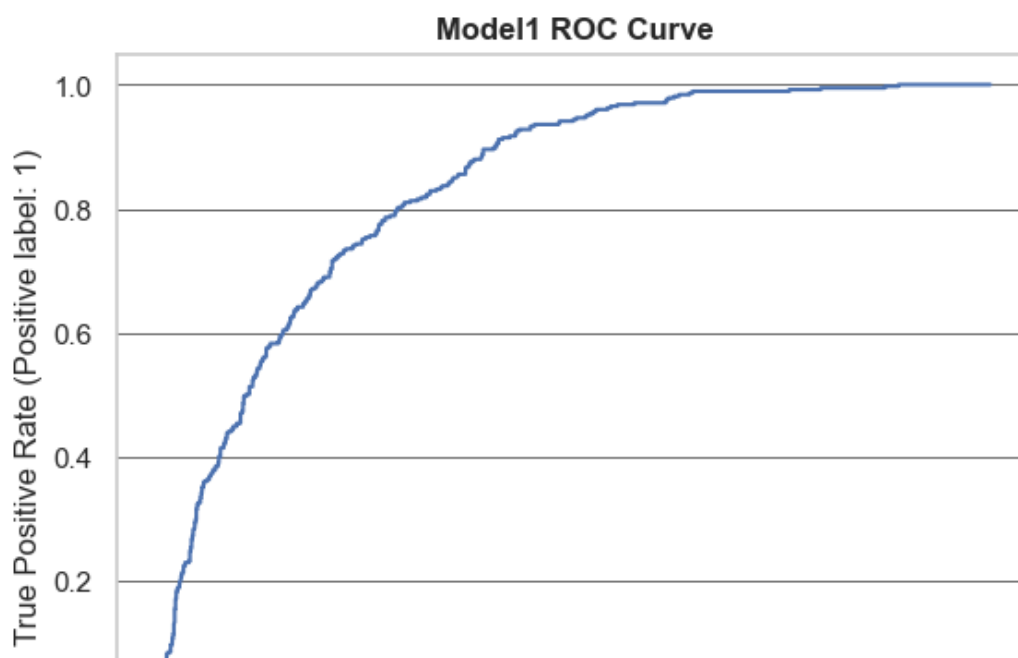
# Model Evaluation
print(f'Model Report: {classification_report(y_test, y_pred2)}')
print(f'\nAccuracy Score: {accuracy_score(y_test, y_pred2)}')
print(f'\nROC_AUC Score: {roc_auc_score(y_test, y_proba2)}')
print(f'\nPrecision Score: {precision_score(y_test, y_pred2)}')
print(f'\nF1 Score: {f1_score(y_test, y_pred2)}')
print(f'\nRecall Score: {recall_score(y_test, y_pred2)}')

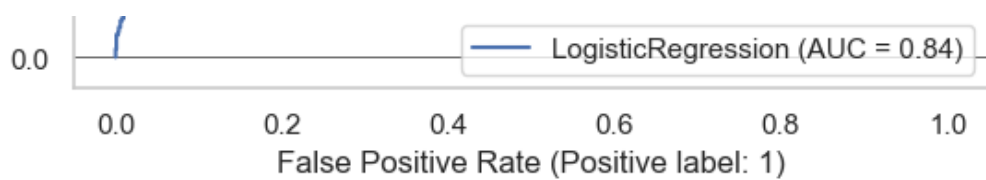
```

<Figure size 800x600 with 0 Axes>



<Figure size 800x600 with 0 Axes>





Model Report:		precision	recall	f1-score	support
	0	0.91	0.71	0.80	1035
	1	0.50	0.80	0.62	374
	accuracy			0.74	1409
	macro avg	0.70	0.76	0.71	1409
	weighted avg	0.80	0.74	0.75	1409

Accuracy Score: 0.7366926898509581

ROC_AUC Score: 0.8353625771784339

Precision Score: 0.5025295109612141

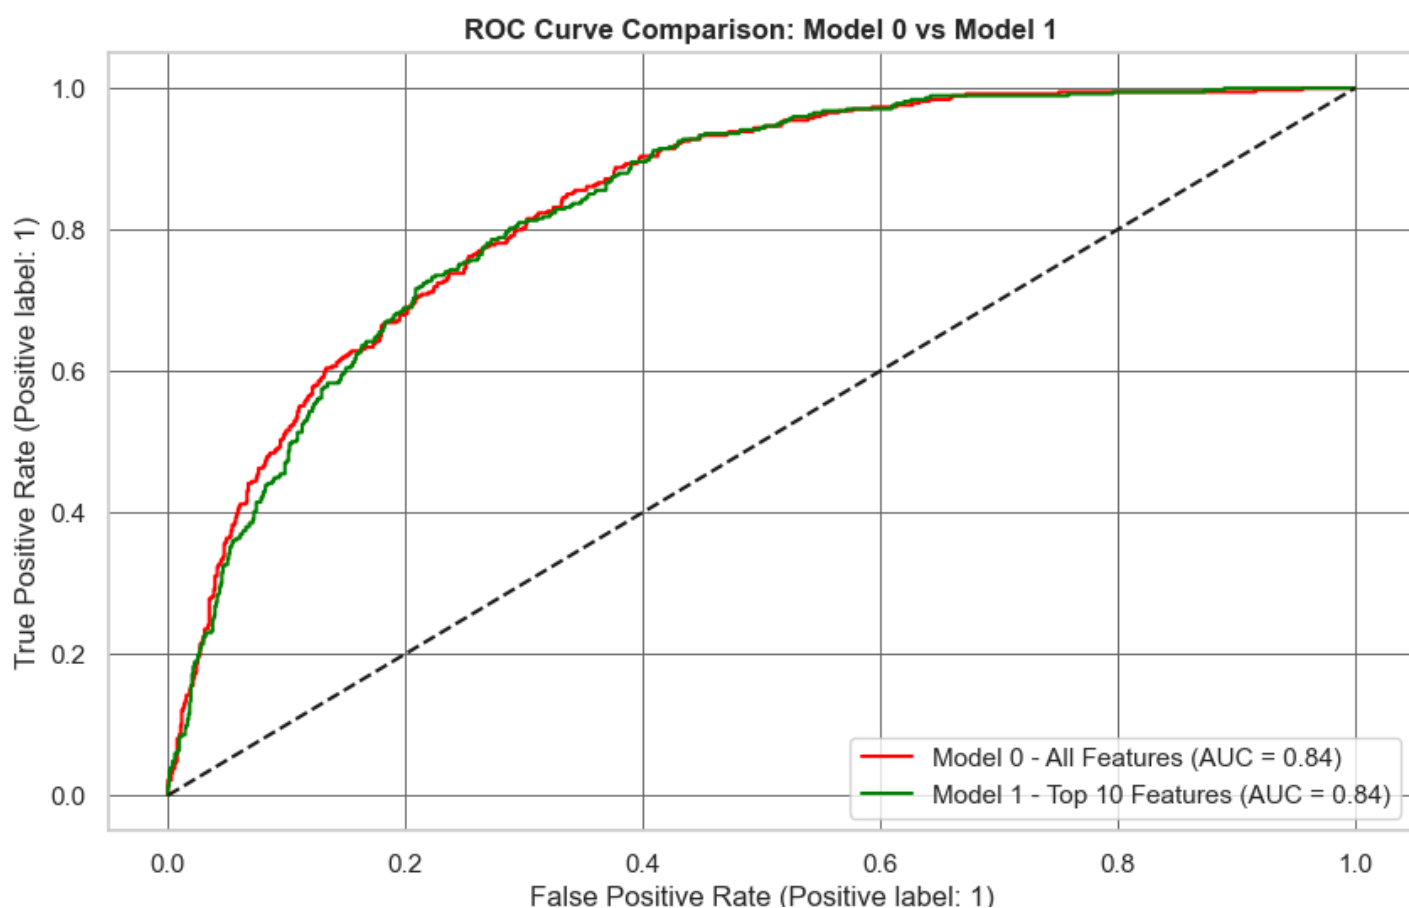
F1 Score: 0.6163391933815925

Recall Score: 0.7967914438502673

Visual comparison of the two models using ROC Curve

In [835]:

```
plt.figure(figsize=(10, 6))
RocCurveDisplay.from_predictions(y_test, y_proba, name='Model 0 - All Features', color='red', ax=plt.gca())
RocCurveDisplay.from_predictions(y_test, y_proba2, name='Model 1 - Top 10 Features', color='green', ax=plt.gca())
plt.title('ROC Curve Comparison: Model 0 vs Model 1', weight='bold')
plt.grid(True)
plt.plot([0, 1], [0, 1], 'k--')
plt.savefig('images\Comparing 2 Logistic models.jpg', dpi=300, bbox_inches='tight')
plt.show()
```



CONCLUSION :

- Both Logistic Regression models (with all features vs. top 10 features) perform almost identically.
- Simplifying the model (Model1) did not compromise performance, which is great.
- But since scores are plateauing, it's time to try a different type of model that can:
 - Capture non-linear relationships,
 - Handle interactions between features,
 - Possibly improve recall or precision, especially on the minority class

Decision Tree Classifier Modeling

Buildin a base Decision Tree Model

In [839]:

```
# Import libraries
from sklearn.tree import DecisionTreeClassifier

# Since the variables were identified in our earlier model we can use them as they are
# Model initialization, fitting and training
dt_model = DecisionTreeClassifier()
dt_model.fit(X_train_scaled, y_train)

# Predict y_test
ydt_pred = dt_model.predict(X_test_scaled)

#Get y probability predictions for roc_auc score
ydt_proba = dt_model.predict_proba(X_test_scaled)[:, 1]

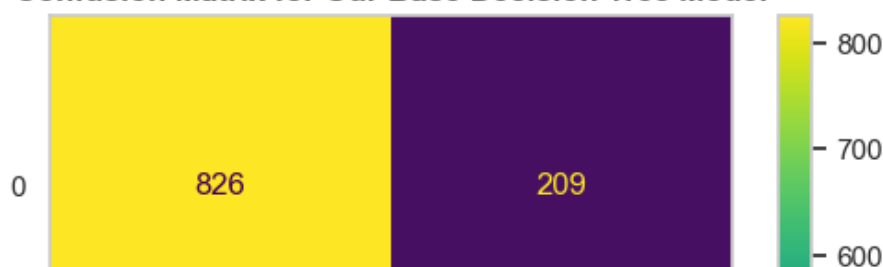
# Plot confusion matrix for our base Decision tree model
dt_cm = confusion_matrix(y_test, ydt_pred)
dt_matrix = ConfusionMatrixDisplay(confusion_matrix = dt_cm)
plt.figure(figsize = (8, 6))
dt_matrix.plot()
plt.title('Confusion Matrix for Our Base Decision Tree Model', weight = 'bold')
plt.grid(False)
plt.show()

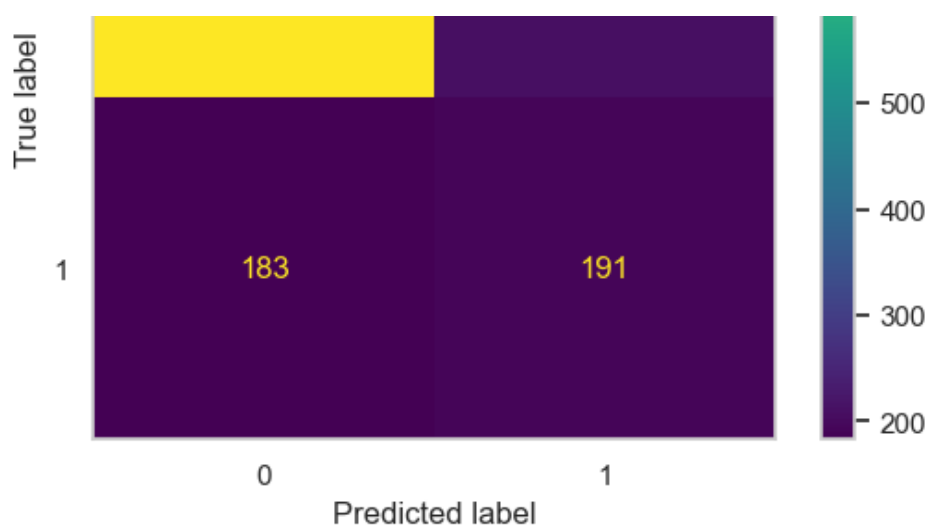
# Ploting ROC curve for our Base Decision Tree Model
plt.figure(figsize = (8, 6))
RocCurveDisplay.from_estimator(dt_model, X_test_scaled, y_test)
plt.title('ROC curve for our Base Decision Tree Model', weight = 'bold')
plt.grid(True)
plt.savefig('images\Base Decision Tree Model.jpg', dpi=300, bbox_inches='tight')
plt.show()

# Model Evaluation Metrics
print(f'Model Report: {classification_report(y_test, ydt_pred)}')
print(f'\nAccuracy Score: {accuracy_score(y_test, ydt_pred)}')
print(f'\nROC_AUC Score: {roc_auc_score(y_test, ydt_proba)}')
print(f'\nPrecision Score: {precision_score(y_test, ydt_pred)}')
print(f'\nF1 Score: {f1_score(y_test, ydt_pred)}')
print(f'\nRecall Score: {recall_score(y_test, ydt_pred)}')
```

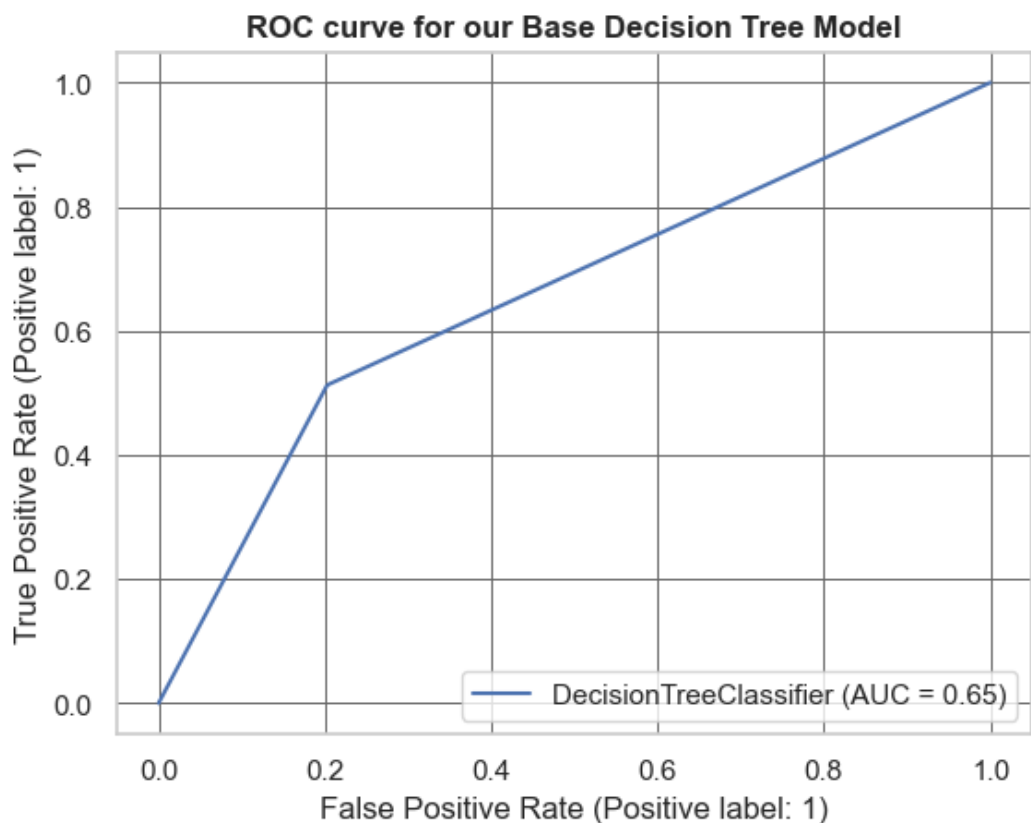
<Figure size 800x600 with 0 Axes>

Confusion Matrix for Our Base Decision Tree Model





<Figure size 800x600 with 0 Axes>



Model Report:		precision	recall	f1-score	support
	0	0.82	0.80	0.81	1035
	1	0.48	0.51	0.49	374
accuracy			0.72	1409	
macro avg		0.65	0.65	0.65	1409
weighted avg		0.73	0.72	0.72	1409

Accuracy Score: 0.7217885024840313

ROC_AUC Score: 0.6549755870727738

Precison Score: 0.4775

F1 Score: 0.49354005167958653

Recall Score: 0.5106951871657754

In [784]:

```
#build fit and train the model with the top 10 featuures
odt_model = DecisionTreeClassifier(class_weight = 'balanced', random_state = 42)
odt_model.fit(X_train_top_scaled, y_train)

# Predict the y_test values
ydt_pred2 = odt_model.predict(X_test_top_scaled)

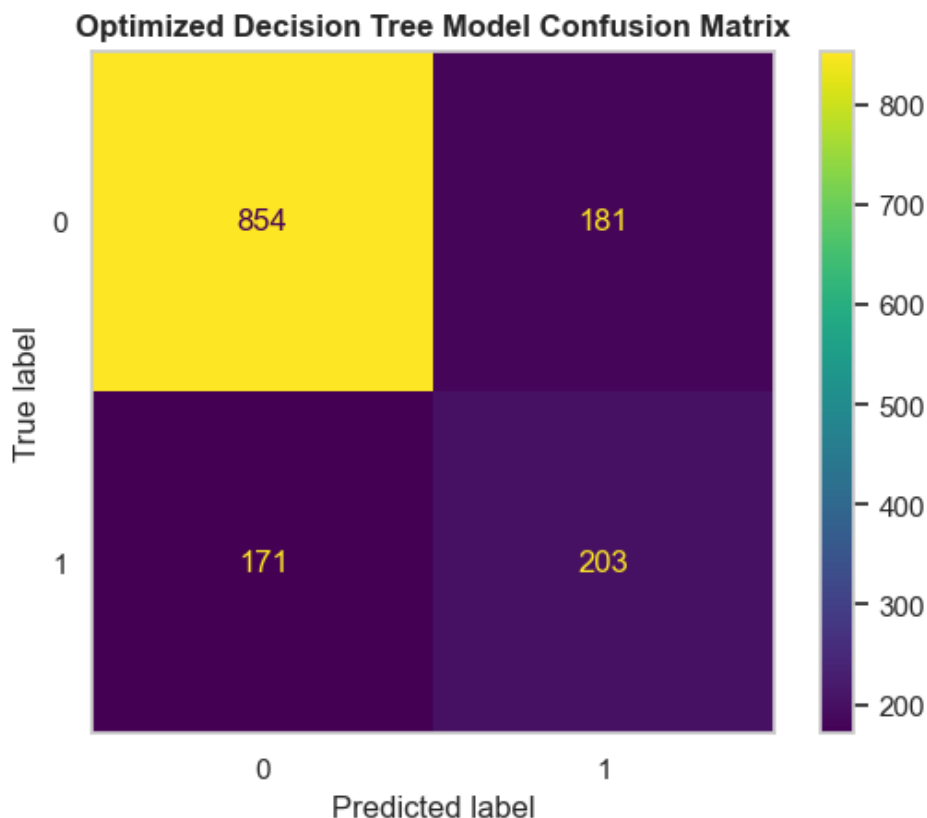
# Get the probability predictions of y_test
ydt_proba2 = odt_model.predict_proba(X_test_top_scaled)[:, 1]

# Plot the confusion matrix for the optimized model
odt_cm = confusion_matrix(y_test, ydt_pred2)
odt_matrix = ConfusionMatrixDisplay(confusion_matrix = odt_cm)
plt.figure(figsize = (8, 6))
odt_matrix.plot()
plt.title('Optimized Decision Tree Model Confusion Matrix', weight = 'bold')
plt.grid(False)
plt.show()

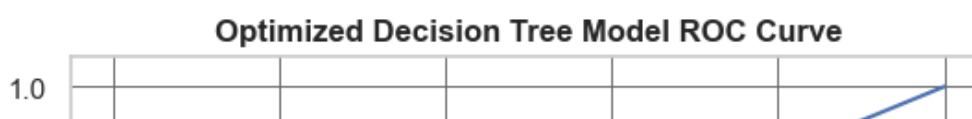
# Plotting ROC Curve for the Optimized Decision Tree Model
plt.figure(figsize = (8, 6))
RocCurveDisplay.from_estimator(odt_model, X_test_top_scaled, y_test)
plt.title('Optimized Decision Tree Model ROC Curve', weight = 'bold')
plt.grid(True)
plt.show()

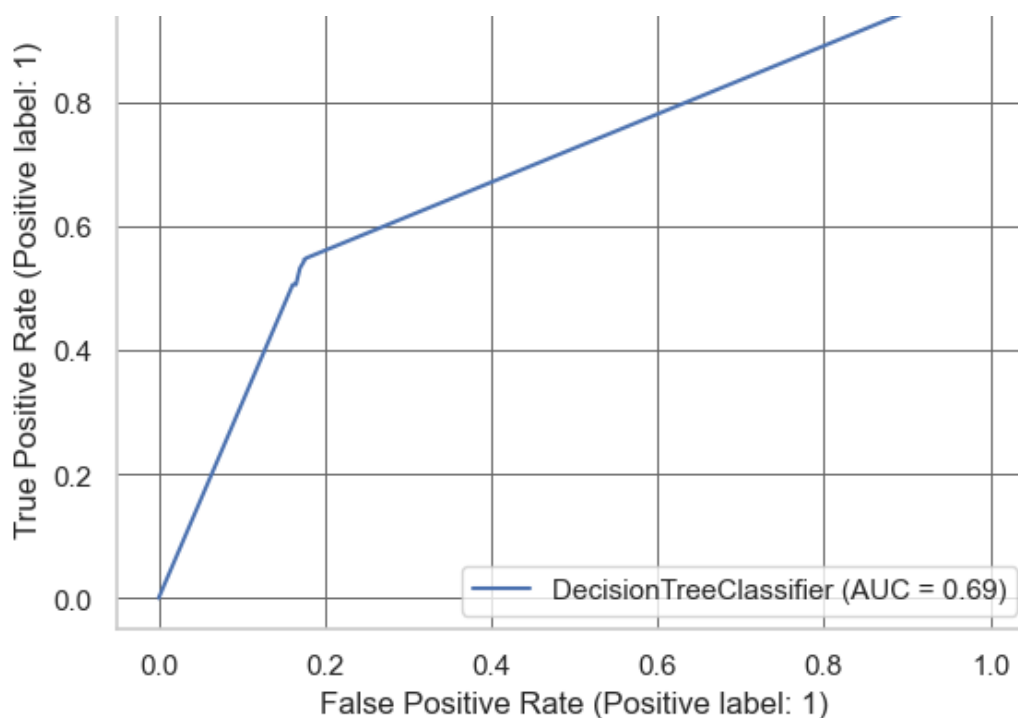
# Evaluating the Optimized Model Metrics
print(f'Model Report: {classification_report(y_test, ydt_pred2)}')
print(f'\nAccuracy Score: {accuracy_score(y_test, ydt_pred2)}')
print(f'\nROC_AUC Score: {roc_auc_score(y_test, ydt_proba2)}')
print(f'\nPrecision Score: {precision_score(y_test, ydt_pred2)}')
print(f'\nF1 Score: {f1_score(y_test, ydt_pred2)}')
print(f'\nRecall Score: {recall_score(y_test, ydt_pred2)}')
```

<Figure size 800x600 with 0 Axes>



<Figure size 800x600 with 0 Axes>





Model Report:		precision	recall	f1-score	support
0	0.83	0.83	0.83	1035	
1	0.53	0.54	0.54	374	
accuracy			0.75	1409	
macro avg		0.68	0.68	0.68	1409
weighted avg		0.75	0.75	0.75	1409

Accuracy Score: 0.7501774308019872

ROC_AUC Score: 0.6858624609263996

Precision Score: 0.5286458333333334

F1 Score: 0.5356200527704486

Recall Score: 0.5427807486631016

Optimizing My Decision Tree Model using GridSearchCV

In [786]:

```
from sklearn.model_selection import GridSearchCV

# Define the model
dt = DecisionTreeClassifier(random_state=42)

# Set the hyperparameter grid
param_grid = {
    'max_depth': [3, 5, 10, 15, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'criterion': ['gini', 'entropy'],
    'class_weight': [None, 'balanced']
}

# Set up GridSearchCV
grid_search = GridSearchCV(
    estimator=dt,
    param_grid=param_grid,
    cv=5,
    scoring='f1',
    n_jobs=-1,
    verbose=1
```

```
)

# Fit the grid search to your training data
grid_search.fit(X_train_top_scaled, y_train)

# Get the best estimator
best_dt = grid_search.best_estimator_
print("Best parameters:", grid_search.best_params_)

# Predict using best model
y_pred_best = best_dt.predict(X_test_top_scaled)
y_proba_best = best_dt.predict_proba(X_test_top_scaled)[:, 1]

# Evaluate
print(classification_report(y_test, y_pred_best))
print("ROC_AUC:", roc_auc_score(y_test, y_proba_best))
```

Fitting 5 folds for each of 180 candidates, totalling 900 fits
 Best parameters: {'class_weight': 'balanced', 'criterion': 'entropy', 'max_depth': 5, 'min_samples_leaf': 2, 'min_samples_split': 2}

	precision	recall	f1-score	support
0	0.89	0.78	0.83	1035
1	0.55	0.73	0.63	374
accuracy			0.77	1409
macro avg	0.72	0.76	0.73	1409
weighted avg	0.80	0.77	0.78	1409

ROC_AUC: 0.8364824717765895

In [787]:

```
# Plotting the confusion and the Roc curve of the GridSearchCV model
# Predict probabilities
y_probs = best_dt.predict_proba(X_test_top_scaled)[:, 1]

# Compute ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_probs)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'AUC = {roc_auc_score(y_test, y_probs):.2f}', color='blue')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.title('ROC Curve - Optimized Decision Tree', weight = 'bold')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

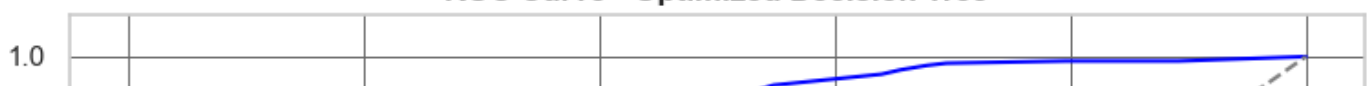
# Confusion Matrix Plot

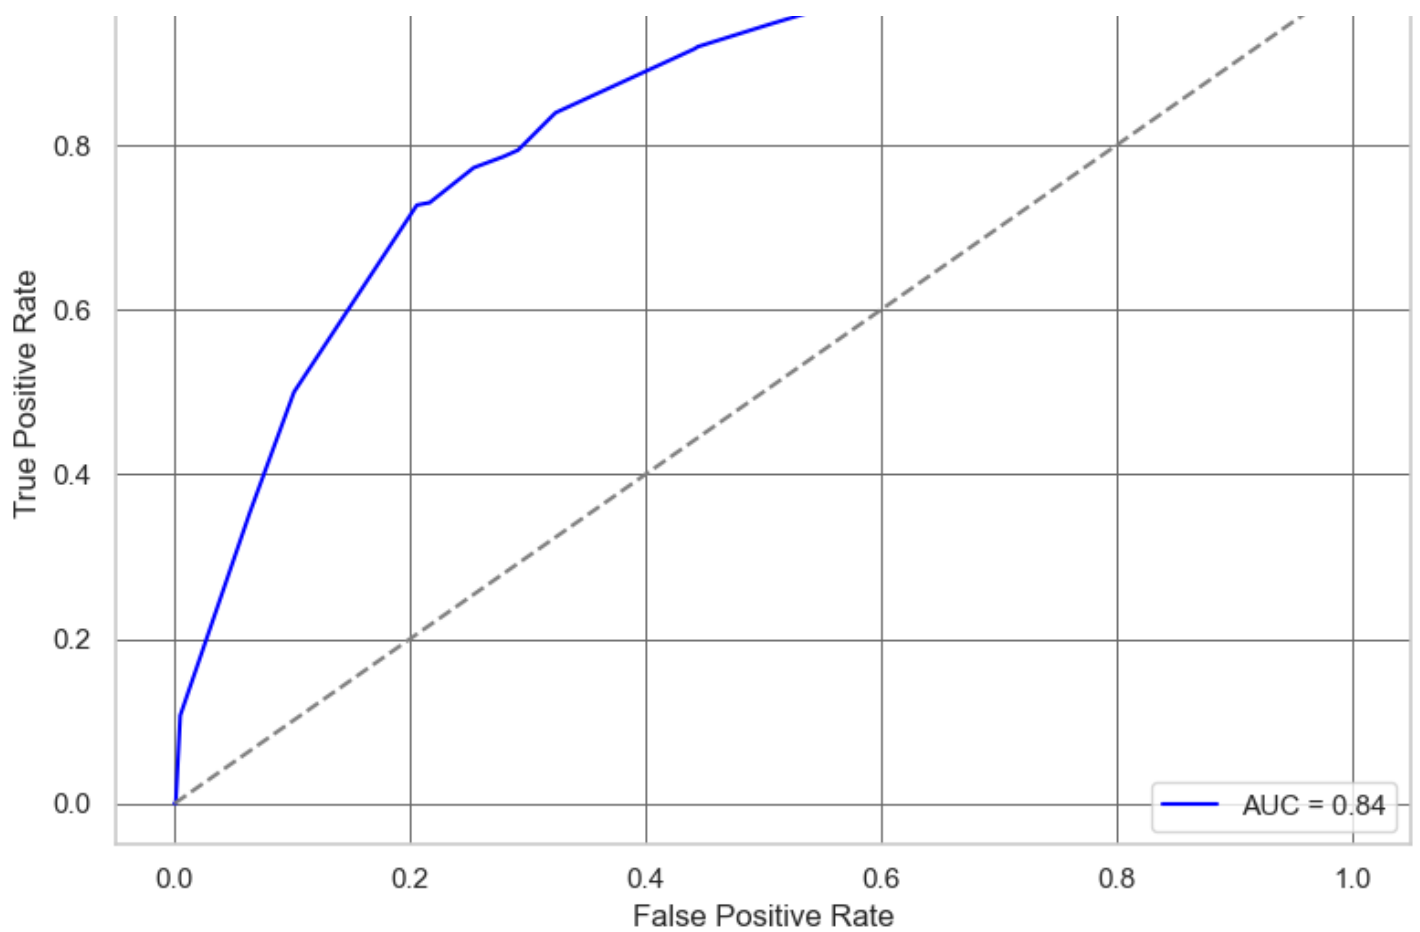
# Predict class labels
y_pred = best_dt.predict(X_test_top_scaled)

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)

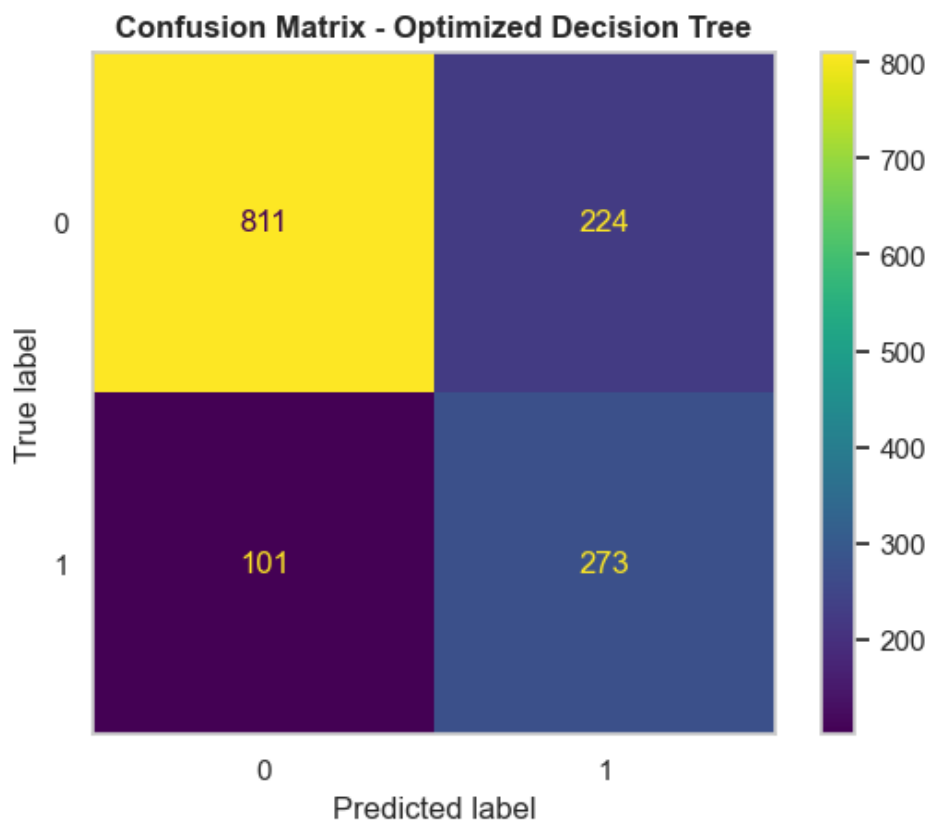
# Plot
plt.figure(figsize = (8, 6))
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=best_dt.classes_)
disp.plot()
plt.title('Confusion Matrix - Optimized Decision Tree', weight = 'bold')
plt.grid(False)
plt.show()
```

ROC Curve - Optimized Decision Tree





<Figure size 800x600 with 0 Axes>



Plotting the tree structure

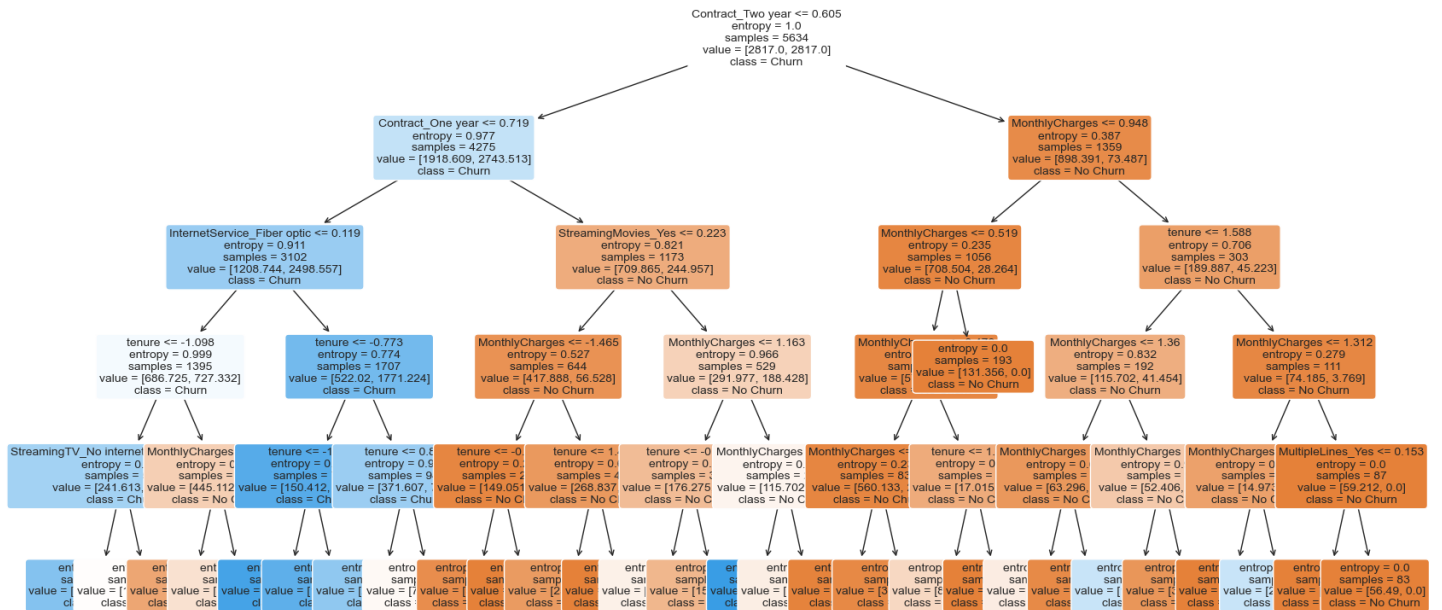
```
In [789]:  
  
# Importing libraries  
from sklearn.tree import plot_tree  
  
plt.figure(figsize=(20, 10)) # Adjust size as needed  
plot_tree(best_dt,
```

```

feature_names=X_train_top.columns, # original feature names before scaling
class_names=['No Churn', 'Churn'],
filled=True,
rounded=True,
fontsize=10)
plt.title("Optimized Decision Tree Structure", weight = 'bold', fontsize = 20)
plt.savefig('images\Decision Tree Structure.jpg', dpi=300, bbox_inches='tight')
plt.show()

```

Optimized Decision Tree Structure



The plot above shows how my decisson tree model is structured

PLoting the 3 Decision Models

In [837]:

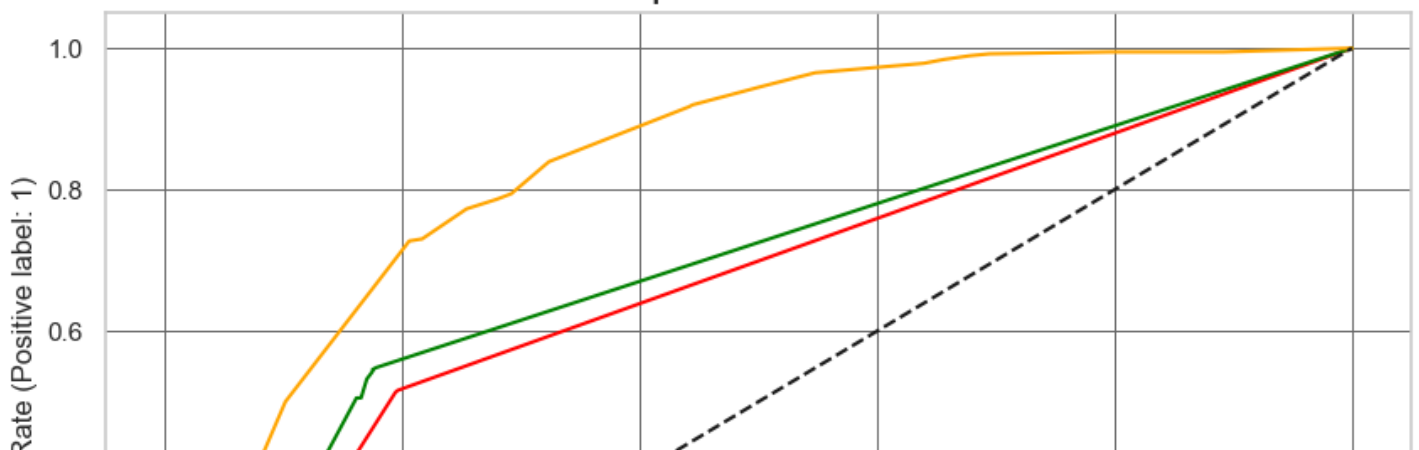
```

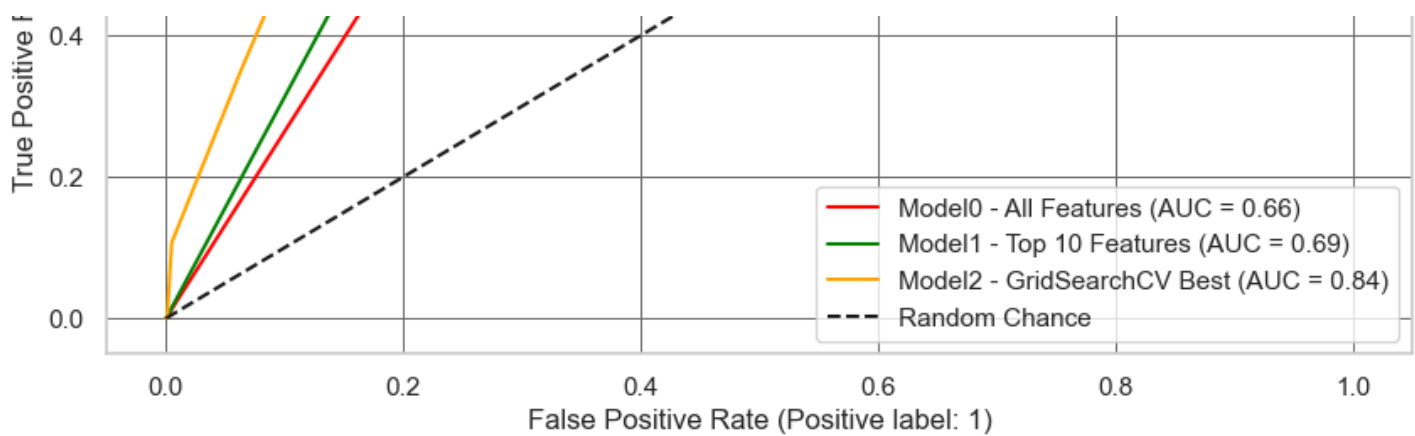
plt.figure(figsize=(10, 6))
RocCurveDisplay.from_predictions(y_test, ydt_proba, name='Model0 - All Features', color=
'red', ax=plt.gca())
RocCurveDisplay.from_predictions(y_test, ydt_proba2, name='Model11 - Top 10 Features', col
or='green', ax=plt.gca())
RocCurveDisplay.from_predictions(y_test, y_proba_best, name='Model2 - GridSearchCV Best'
, color='orange', ax=plt.gca())

plt.plot([0, 1], [0, 1], 'k--', label='Random Chance')
plt.title('ROC Curve Comparison of Decision Tree Models', weight='bold')
plt.grid(True)
plt.legend()
plt.savefig('images\Comparing 3 diffrent decisson tree models.jpg', dpi=300, bbox_inches
='tight')
plt.show()

```

ROC Curve Comparison of Decision Tree Models





CONCLUSION

- Using GridSearchCV greatly improved the performance of my model
- Compared to the base model and the top 10 feature model:
 - The ROC-AUC score improved
 - Better recall for the minority class
 - More balanced precision and f1-score
- This highlights the importance of hyperparameter tuning in boosting classification model performance, especially in imbalanced datasets.

Conclusion

In this project, we successfully built and evaluated several classification models to predict customer churn in a telecom company. Through thorough data cleaning, exploration, and modeling, we identified patterns and customer behaviors associated with churn.

Key insights include:

- Contract type plays a major role: customers on month-to-month contracts are significantly more likely to churn than those with one or two-year contracts.
- Lack of tech support and security services correlates strongly with churn.
- Electronic payment methods show a higher churn rate compared to other payment methods.
- Customers with short tenure, indicating newer users, are more likely to leave.

Modeling results:

- Our baseline Decision Tree model achieved fair performance, but using top features and applying hyperparameter tuning via GridSearchCV significantly improved metrics.
- The best performing model (optimized Decision Tree) achieved:
 - Accuracy: ~77%
 - Recall (Churn class): ~73%
 - ROC AUC: ~0.84
- This model is interpretable, aligns well with business logic, and provides insight into which features drive churn decisions.

Recommendations

Based on our findings, we recommend the telecom provider take the following actions to reduce churn:

1. Incentivize Long-Term Contracts

- Offer discounts or loyalty rewards for 1- or 2-year contracts to discourage month-to-month churners.

2. Improve Tech Support and Security Offerings

- Promote bundled services including Tech Support, Online Backup, and Device Protection.
- Ensure customers are aware of these services and their benefits.

3. Address Short-Tenure Churn

- Implement an onboarding and retention strategy within the first 6 months to engage new customers and reduce early churn.

4. Reconsider Payment Method Incentives

- Encourage alternative payment methods over electronic checks, which correlate with higher churn.
- Explore if this is related to demographic or service dissatisfaction.

5. Target At-Risk Segments

- Use the model to score and flag high-risk customers regularly.
- Send targeted retention offers or interventions to those most likely to churn.

STREAMLIT DEMO

In [849]:

```
import pickle

with open('best_model.pkl', 'wb') as file:
    pickle.dump(best_dt, file)

with open('scaler.pkl', 'wb') as f:
    pickle.dump(scaler1, f)
```

In []: